

A2A Security Framework (Agents→Automation)

End-to-end, LLM-assisted SOC orchestration scaffold built around **LangGraph** + **Pydantic** types and a clean separation of concerns (domains, agents, tools, and graph). This README walks through the repository structure and documents **every script and function** so you can customize it to your stack (SIEM/EDR/CMDB, XSOAR, Slack/PagerDuty, etc.).

Quick Start

Prerequisites

- **Python** 3.11+ (3.12 recommended)
- (For the Streamlit UI) **Graphviz** runtime
 - macOS: `brew install graphviz`
 - Ubuntu/Debian: `sudo apt-get install graphviz`
 - Windows (choco): `choco install graphviz`
- **OpenAI API key** (or swap in your LLM provider in `src/agents.py`)

Setup

```
git clone <your-repo-url>
cd security_framework_a2a
python -m venv .venv
source .venv/bin/activate  # Windows: .venv\Scripts\activate
pip install -r requirements.txt
```

Create a `.env` with your secrets (at minimum the LLM key):

```
# .env
OPENAI_API_KEY=sk-...
# Optional overrides
A2A_POLICY=./config.yaml          # path to policy file
A2A_OBS_LOG=./a2a_observability.log
```

Run the CLI demo

```
python src/run_demo.py
```

This loads a sample `Alert`, runs the graph once, and prints the **decision**, **status**, **chosen playbooks**, and an **audit log**.

Run the Streamlit UI

```
streamlit run src/streamlit_app.py
```

- Choose **Form** or **JSON** to input an alert - Pick **Native streaming (parallel)** or **Sequential stepper** - Watch node statuses, metrics, enrichment, and timeline update live - Uses an on-disk SQLite checkpointer if available (falls back to memory)

Repository Layout

```
security_framework_a2a/
  .env
  .gitignore
  a2a_observability.log
  config.yaml
  requirements.txt
  src/
    agents.py
    config.py
    domains.py
    graph.py
    run_demo.py
    streamlit_app.py
    tools.py
```

- `config.yaml` — Policy thresholds & future playbook rules
- `a2a_observability.log` — JSONL events emitted during runs
- `requirements.txt` — Pinned Python deps (LangChain, LangGraph, Streamlit, etc.)
- `src/` — All code: typed domains, LLM agents, orchestration graph, UI, and integration stubs

Overall Flow (Graph Topology)

High-level pipeline (left→right), with the final decision branching to automation or human triage:

```
graph LR
  A[enrich] --> B[validity]
  B --> C[severity]
  C --> D[exploitability]
  D --> E[playbooks]
  E --> F[decision]
  F -->|UPLOAD_XSOAR| G[upload_xsoar]
```

```
F -->|SOC_TRIAGE| H[soc_triage]  
G --> I[update_status]  
H --> I[update_status]
```

The Streamlit app renders a Graphviz version of the same topology and streams node-level updates.

src/domains.py — Typed Domain Objects

Strict, LLM-friendly Pydantic models that define the **contract** between agents and the graph.

class Indicator

- **type**: one of ip | domain | hash | url | email | user | process | file | host
- **value**: indicator string
- **context**: optional human hint (e.g., "tor exit", "prod host")

class Alert

- **id**: unique alert id (e.g., ALRT-1001)
- **source**: default "SIEM"
- **title, description**: human description
- **indicators**: list of Indicator
- **created_at**: timestamp (default now)

class ValidityScore

- **label**: "False Positive" | "False Negative" | "True Positive" | "True Negative"
- **likelihood**: 0..1 belief that alert is real (TP)
- **rationale**: brief explanation

class SeverityScore

- **level**: 1 | 2 | 3 (map to your program's levels)
- **impact**: "Low" | "Medium" | "High"
- **rationale**

class ExploitabilityScore

- **level**: "Not Exploitable" | "Low" | "Medium" | "Critical"
- **likelihood**: 0..1 belief the condition can be exploited
- **rationale**

class PlaybookChoice

- **names**: list of playbook names to run (e.g., Contain_Endpoint , Reset_Creds)
- **rationale**

class Decision

- **soc_attention**: boolean gate telling if the incident deserves automation
- **path**: "UPLOAD_XSOAR" | "SOC_TRIAGE"
- **rationale**: short justification

class ActionLog

- **at**: timestamp
- **event**: event key (e.g., enriched, xsoar_uploaded)
- **details**: JSON blob with any fields you add (host, counts, api responses)

class SOCState (TypedDict) — Shared Graph State

- **alert**: Alert (*input*)
- **enrichment**: dict (*output of* enrich_node)
- **validity**: ValidityScore
- **severity**: SeverityScore
- **exploitability**: ExploitabilityScore
- **playbooks**: PlaybookChoice
- **decision**: Decision
- **status**: str (human-readable status summary)
- **logs**: list[ActionLog] — reducer is operator .add to support parallel node appends

These types are the **single source of truth** for all LLM structured outputs and node IO.

src/config.py & config.yaml — Policy Loader & Thresholds

- `load_policy()` loads YAML from A2A_POLICY env var or the repo-root config.yaml, merges onto internal defaults:
- `thresholds.validity_tp_min` (default **0.60** unless overridden by config.yaml)
- `thresholds.severity_min` (default **2**)
- `thresholds.exploit_levels_escalate` (default **["Medium", "Critical"]**)
- `future: playbook_rules` mapping

Example config.yaml (in repo):

```
thresholds:  
    validity_tp_min: 0.65  
    severity_min: 2  
    exploit_levels_escalate: ["Medium", "Critical"]  
    # playbook_rules: []
```

src/agents.py — LLM Agents

Each agent uses a shared LLM client and returns a **typed** object from `domains.py`.

get_llm()

Returns a preconfigured `ChatOpenAI` client (`model="gpt-4o-mini"`, `temperature=0`). Swap to your provider if needed.

validity_agent(alert: Alert, enrichment: dict) -> ValidityScore

- Prompt asks the model to determine whether the alert is real
- Returns **label**, **likelihood**, and **rationale** as a `ValidityScore`

severity_agent(alert: Alert, enrichment: dict) -> SeverityScore

- Scores **business/technical impact** (1-3 & Low/Medium/High) with a short justification

exploitability_agent(alert: Alert, enrichment: dict) -> ExploitabilityScore

- Estimates exploitability and likelihood (0-1) with rationale

playbook_agent(alert: Alert, enrichment: dict, scores: ...) -> PlaybookChoice

- Chooses playbooks to run based on observed context & scores (extensible with future `playbook_rules`)

decision_agent(validity, severity, exploitability, policy=None) -> Decision

- Applies a **deterministic gate** using thresholds from `load_policy()`
- Escalate when `validity.likelihood ≥ validity_tp_min` and (`severity.level ≥ severity_min` or `exploitability.level ∈ exploit_levels_escalate`)
- Calls the LLM to produce a concise **rationale** and **path** (`UPLOAD_XSOAR` vs `SOC_TRIAGE`)

All agent prompts are created with `ChatPromptTemplate` and use `with_structured_output(...)` to guarantee schema.

src/tools.py — Integration Stubs

Replace these stubs with your real systems.

```
search_internal_data(query: str) -> list[dict]
```

Call your SIEM/EDR/CMDB/ticketing/asset DB. Stub returns example EDR/CMDB hits.

```
search_external_osint(query: str) -> list[dict]
```

Hook to VirusTotal, GreyNoise, AbuseIPDB, etc. Stub returns canned examples.

```
@dataclass class XSOARClient
```

- `upload_incident(payload: dict) -> dict` — POST to your XSOAR/Demisto
- `update_status(incident_id: str, status: str) -> dict`

```
notify_soc(payload: dict) -> dict
```

Send to your human triage channel: Slack, PagerDuty, email, etc.

src/graph.py — Orchestration (LangGraph)

Defines graph **nodes**, **edges**, observability, and build helpers.

Observability

- `_emit_observable(event, **details)` appends JSON lines to `A2A_OBS_LOG` (default `a2a_observability.log`)
- `log_event(event, **details) -> list[ActionLog]` produces a delta list for the **reducer** to append
- `log(state, event, **details) -> list[ActionLog]` also emits to the file and returns the updated list

Node implementations

- `enrich_node(state) -> state_delta`
- For each `Indicator`, queries `search_internal_data` and OSINT for net IOCs
- Returns `{ "enrichment": {internal, external}, "logs": [...] }`
- `validity_node(state) -> state_delta`
- Calls `validity_agent`; returns `{ "validity": ValidityScore, "logs": [...] }`
- `severity_node(state) -> state_delta`
- Calls `severity_agent`; returns `{ "severity": SeverityScore, "logs": [...] }`

- `exploitability_node(state) -> state_delta`
- Calls `exploitability_agent`; returns `{ "exploitability": ExploitabilityScore, "logs": [...] }`
- `playbooks_node(state) -> state_delta`
- Calls `playbook_agent`; returns `{ "playbooks": PlaybookChoice, "logs": [...] }`
- `decision_node(state) -> state_delta`
- Calls `decision_agent`; returns `{ "decision": Decision, "logs": [...] }`
- `upload_xsoar_node(state) -> state_delta`
- Builds a minimal incident payload and calls `XSOARClient.upload_incident`
- Returns status and `xsoar_uploaded` log
- `soc_triage_node(state) -> state_delta`
- Notifies human queue via `notify_soc`
- Returns status and `triage_notified` log
- `update_status_node(state) -> state_delta`
- If path was `UPLOAD_XSOAR`, updates status via `XSOARClient.update_status`
- Else no-op, logs `status_noop`

Build & topology helpers

- `build_graph(parallel: bool = True, checkpointer=None)`
- Registers nodes, wires edges, configures START/END, and compiles the app
- Pass a `SqliteSaver` / `MemorySaver` to enable checkpointing & streaming
- `describe_topology()` returns `{ nodes, edges }` for visualization

`src/run_demo.py` — Minimal CLI Runner

- Loads `.env`
- Creates an example `Alert` (ip/host/user)
- `app = build_graph() then app.invoke({"alert": alert, "logs": []})`
- Prints the resulting `decision`, `status`, `playbooks`, and the timeline of `logs`

src/streamlit_app.py — Interactive UI

A compact UI to explore the graph live.

Key pieces

- **Graphviz diagram** with node coloring based on status (pending/running/done/branch)
- **Two run modes**
- **Native streaming (parallel)**: uses LangGraph streaming to update the UI as deltas arrive
- **Sequential stepper**: runs each node with a user-controlled delay — helpful for demoing
- **Checkpointing**
 - Tries `langgraph.checkpoint.sqlite.SqliteSaver("a2a_checkpoints.sqlite")`
 - Falls back to `MemorySaver` if SQLite is unavailable
- Enter a **Thread ID** to reuse/inspect a previous run's state
- **Sidebar**
 - Input mode: **Form** or **JSON** (raw `Alert` `JSON`)
 - Policy viewer (current thresholds from `config.yaml`)
- **Panels**
 - Metrics: validity/severity/exploitability/playbooks/decision
 - Enrichment (internal & external hits)
 - Timeline of `ActionLog` events
 - Final state + downloadable JSON snapshot

Configuration & Environment

- **LLM**: controlled in `agents.get_llm()` — swap provider/model or add retries/timeouts
- **Policy**: edit `config.yaml` or set `A2A_POLICY=/path/to/your.yaml`
- **Observability**: set `A2A_OBS_LOG=/path/to/file.jsonl`
- **Secrets**: put keys in `.env` (dotenv is loaded in CLI & UI)

Extending the Framework

1. **Wire real data sources** in `tools.py` (SIEM/EDR/CMDB/ticketing/OSINT)
2. **Harden prompts** or add **guardrails** (schema validators, unit tests)
3. **Add rules** (`playbook_rules` in `config.yaml`) to enforce hard decisions
4. **Introduce new nodes** (e.g., `contain`, `notify_owner`) — add in `graph.py`, then include in the UI's `NODES`
5. **Swap the LLM** (e.g., on-prem, Azure OpenAI, Bedrock, etc.)

Example: Alert JSON

```
{  
  "id": "ALRT-1001",  
  "title": "Outbound connection to known brute-force IP",  
  "description": "FW logs show repeated egress to suspicious IP by host  
  srv-42.",  
  "indicators": [  
    { "type": "ip", "value": "203.0.113.55" },  
    { "type": "host", "value": "srv-42" },  
    { "type": "user", "value": "svc-backup" }  
  ]  
}
```

Troubleshooting

- **Graphviz not found** → install it and restart Streamlit (see prerequisites)
- **No streaming/graph errors** → ensure LangGraph version matches `requirements.txt`
- **LLM auth errors** → check `OPENAI_API_KEY` in `.env`
- **Nothing written to** `a2a_observability.log` → confirm `A2A_OBS_LOG` and file permissions

License

Add your license of choice (e.g., MIT) at the repo root.

Attributions

- Built on **LangGraph/LangChain, Streamlit, Pydantic**.